




Build Data Pipelines with Lakeflow Declarative Pipelines



Databricks Academy



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Course Learning Objectives

- Understand the core concepts and components of Lakeflow Declarative Pipelines, including the function and differences between streaming tables, materialized views, and temporary views.
- Identify and configure Lakeflow pipeline settings, such as compute, data assets, trigger modes, and advanced options.
- Develop a functional Lakeflow Declarative Pipeline using the new pipeline editor and SQL-based syntax.
- Incorporate data quality expectations into a pipeline to validate and enforce data integrity.
- Analyze event logs and pipeline metrics to understand the full execution and lifecycle of a Lakeflow Declarative Pipeline.
- Design and implement a Change Data Capture (CDC) to a pipeline using APPLY CHANGES INTO to handle slowly changing dimensions (SCD).



Agenda

Course Sections

- **Introduction to Data Engineering in Databricks**
- **Lakeflow Declarative Pipeline Fundamentals**
- **Building Lakeflow Declarative Pipelines**



Course Prerequisites (REQUIRED)



Fundamental Knowledge of the Databricks Platform

- Databricks Workspaces
- Apache Spark
- Delta Lake and the Medallion Architecture
- Unity Catalog



Course Prerequisites (REQUIRED)



Fundamental Knowledge
of the Databricks Platform

- Databricks Workspaces
- Apache Spark
- Delta Lake and the Medallion Architecture
- Unity Catalog



**Knowledge of Ingesting
Raw Data into a Table**

- Familiarity with ingesting raw data source files using the **read_files** SQL function
 - CSV
 - JSON
 - TXT
 - Parquet



Course Prerequisites (REQUIRED)



Fundamental Knowledge of the Databricks Platform

- Databricks Workspaces
- Apache Spark
- Delta Lake and the Medallion Architecture
- Unity Catalog



Knowledge of Ingesting Raw Data into a Table

- Familiarity with ingesting raw data source files using the `read_files` SQL function
 - CSV
 - JSON
 - TXT
 - Parquet



Experience Transforming Data with SQL

- Experience in writing **intermediate-level queries** using SQL
- Basic knowledge of **SQL Joins**



Lab Exercise Environment




Technical Details

- Your lab environment is provided by Vocareum.
- It will open in a new tab.
- It has been configured with the permissions and resources required to accomplish the tasks outlined in the lab exercise.
- Third party cookies must be enabled in your browser for Vocareum's user experience to work properly.
- Make sure to enable pop ups!





Introduction to Data Engineering in Databricks



Build Data Pipelines with Lakeflow Declarative Pipelines



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Section Learning Objectives

- Understand the components of Data Engineering with Lakeflow in Databricks with Connect, Declarative Pipelines, and Jobs.
- Explain how Lakeflow Declarative Pipelines incrementally processes data using Streaming Tables and Materialized Views in batch or streaming workloads.
- Demonstrate how to create a Lakeflow Declarative Pipelines project in a Databricks Workspace.



Agenda

Section Overview – Introduction to Data Engineering in Databricks

- **Data Engineering in Databricks**
- **What are Lakeflow Declarative Pipelines?**
- **Course Setup and Creating a Pipeline**
- **Course Project Overview**





Introduction to Data Engineering in Databricks

LECTURE

Data Engineering in Databricks



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Data Engineering in Databricks



databricks **LAKEFLOW**

UNIFIED DATA ENGINEERING FOR THE DATA INTELLIGENCE PLATFORM

CONNECT



Efficient
ingestion
connectors

DECLARATIVE PIPELINES (Previously DLT)



Accelerated ETL
development

JOBS (Previously Workflows)



Reliable
orchestration for
analytics and AI

INDUSTRY LEADING DATA PROCESSING ENGINE

(Apache Spark + Structured Streaming)

UNIFIED GOVERNANCE

 **Unity Catalog**

OPTIMIZED STORAGE

 **DELTA LAKE**  **Parquet**  **ICEBERG**



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Data Engineering in Databricks

Simplifying batch and streaming ETL with automated reliability, optimizations and built-in data quality with **Lakeflow Declarative Pipelines!**




databricks

LAKEFLOW

UNIFIED DATA ENGINEERING FOR THE DATA INTELLIGENCE PLATFORM

CONNECT



Efficient
ingestion
connectors

DECLARATIVE PIPELINES
(Previously DLT)



Accelerated ETL
development

JOBS

(Previously Workflows)



Reliable
orchestration for
analytics and AI

INDUSTRY LEADING DATA PROCESSING ENGINE
(Apache Spark + Structured Streaming)

UNIFIED GOVERNANCE

 **Unity Catalog**

OPTIMIZED STORAGE

 **DELTA LAKE**

 **Parquet**  **ICEBERG**



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Introduction to Data Engineering in Databricks


LECTURE

What are Lakeflow Declarative Pipelines?




What are Lakeflow Declarative Pipelines?

It's difficult to build and operate reliable data pipelines




What are Lakeflow Declarative Pipelines?

It's difficult to build and operate reliable data pipelines



What are Lakeflow Declarative Pipelines?

Lakeflow Declarative Pipelines: Reliable data pipelines made easy



What are Lakeflow Declarative Pipelines?

Introducing Declarative Pipelines: Reliable data pipelines made easy

- **Simplified Pipeline Authoring**


Easily declare data ingestion and transformations with **SQL or Python**, and let Lakeflow Declarative Pipelines handle the rest!

- **Intelligent Optimization at Scale**

Automated scaling and recovery improve reliability and reduce maintenance.


- **Unified Batch and Streaming**

Pipelines seamlessly adapt to **both near real-time and batch workloads**, optimizing performance and cost.



What are Lakeflow Declarative Pipelines?


Connecting to Data Sources



What are Lakeflow Declarative Pipelines?


Simplifying Batch and Streaming ETL in the Medallion Architecture

Lakeflow Declarative Pipelines




What are Lakeflow Declarative Pipelines?

Overview of Incremental Processing in Declarative Pipelines




What are Lakeflow Declarative Pipelines?

Overview of Incremental Processing in Declarative Pipelines (Run 1)




What are Lakeflow Declarative Pipelines?

Overview of Incremental Processing in Declarative Pipelines (Run 2)




Creating a Pipeline

1



2






Introduction to Data Engineering in Databricks

DEMONSTRATION

REQUIRED - Course Setup and Creating a Pipeline



Notebook: 1 – REQUIRED – Course Setup and Creating a Pipeline



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Introduction to Data Engineering in Databricks


LECTURE

Course Project Overview



Course Project Overview

Cloud Storage





Lakeflow Declarative Pipeline Fundamentals



Build Data Pipelines with Lakeflow Declarative Pipelines



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Section Learning Objectives

- Understand the core concepts and components of Lakeflow Declarative Pipelines pipelines, including how streaming tables, materialized views, and temporary views function and differ.
- Identify and configure pipeline settings such as compute, data assets, trigger modes, and advanced options.
- Develop a functional Declarative Pipeline using the new pipeline editor and SQL-based syntax for the pipeline.
- Incorporate data quality expectations into a Declarative Pipeline pipeline to validate and enforce data integrity.



Agenda

Section Overview – Pipeline Fundamentals

- **Dataset Types Overview**
- **Simplified Pipeline Development**
- **Common Pipeline Settings**
- **Developing a Simple Pipeline**
- **Ensure Data Quality with Expectations**





Lakeflow Declarative Pipeline Fundamentals

LECTURE

Dataset Types Overview



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Dataset Types Overview

How streaming tables, materialized views, and views process data



STREAMING TABLE (ST)

A table with support for streaming or incremental data processing, **only processing new data**



MATERIALIZED VIEW (MV)



VIEWS



Dataset Types Overview

STREAMING TABLE Overview

STREAMING TABLE

- Support for **batch** or **streaming** for **incremental** data processing (exactly **once processing** of the files)
- Each time an streaming table is refreshed, data added to the source tables is **appended** to the **streaming table**
- SQL code that creates streaming tables:
CREATE OR REFRESH STREAMING TABLE
- In SQL, you must use **FROM STREAM read_files()** to **invoke Auto Loader functionality** for incremental streaming reads and checkpoints



Dataset Types Overview

Create a STREAMING TABLE with SQL

Create a STREAMING TABLE named **orders_bronze** from JSON files

```
CREATE OR REFRESH STREAMING TABLE 1_bronze_db.orders_bronze AS
```

```
SELECT
```

```
* ,  
current_timestamp() AS processing_time,  
_metadata.file_name AS source_file
```

```
FROM STREAM read_files(
```

```
  "/Volumes/dbacademy/ops/labuser/orders",  
  format => 'JSON');
```

Creates a streaming table named **orders_bronze** in the **1_bronze_db** schema

Use the **SELECT** clause to specify the columns to read into the streaming table from your data source

Note: The query is using the default catalog

- The **STREAM** keyword indicates to use streaming read for the source files (JSON)
- The **read_files()** function reads files under a provided location and returns the data in tabular form



Dataset Types Overview

Create a STREAMING TABLE with SQL

Create a STREAMING TABLE named **orders_silver** that reads from the **STREAMING orders_bronze** table

```
CREATE OR REFRESH STREAMING TABLE 2_silver_db.orders_silver AS
```

```
SELECT
```

```
    order_id,  
    timestamp(order_timestamp) AS order_timestamp,  
    customer_id,  
    notifications
```

```
FROM STREAM 1_bronze_db.orders_bronze;
```

Specify your SQL transformations for the streaming data from the bronze table

Use STREAM keyword with the name of the STREAMING TABLE you are reading from to transform new rows

Note: The query is using the default catalog



Dataset Types Overview

How streaming tables, materialized views, and views process data



STREAMING TABLE (ST)

A table with support for streaming or incremental data processing, **only processing new data**



MATERIALIZED VIEW (MV)

- Records are processed as required to return accurate results for the **current data state**
- Used for data processing tasks such as:
 - Transformations
 - Aggregations
 - pre-computing slow queries
 - frequently used computations



VIEWS




Dataset Types Overview

MATERIALIZED VIEW Overview

MATERIALIZED VIEW

- Each time a materialized view is updated, **query results are recalculated to reflect changes in upstream datasets**
- Created and **kept up-to-date** by pipeline
- Use the **CREATE OR REFRESH MATERIALIZED VIEW** syntax
- Can be used **anywhere in your pipeline**, not just in the gold layer
- **Where applicable**, results are **incrementally refreshed for materialized views**, avoiding the need to completely rebuild the materialized view when new data arrives (**Serverless Only**)
- **Incremental refresh for materialized views** is a cost-based optimizer to power fast and efficient transformations for materialized views on Serverless compute



Dataset Types Overview

Create a MATERIALIZED VIEW with SQL

Create a MATERIALIZED VIEW named **gold_orders_by_date** that summarizes the **orders_silver** table

```
CREATE OR REFRESH MATERIALIZED VIEW 3_gold_db.gold_orders_by_date AS  
SELECT  
    date(order_timestamp) AS order_date,  
    count(*) AS total_daily_orders  
FROM 2_silver_db.orders_silver  
GROUP BY date(order_timestamp);
```

Create the materialized view **gold_orders_by_date** in the **3_gold_db** schema (database)

Specify the the source table without the STREAM keyword to utilize the source STREAMING table

NOTE: Where possible, queries will be incrementally refreshed

Note: The query is using the default catalog



Dataset Types Overview

How streaming tables, materialized views, and views process data



STREAMING TABLE (ST)

A table with support for streaming or incremental data processing, **only processing new data**



MATERIALIZED VIEW (MV)

- Records are processed as required to return accurate results for the **current data state**
- Used for data processing tasks such:
 - Transformations
 - Aggregations
 - pre-computing slow queries
 - frequently used computations.



VIEWS

Constructs a **virtual table** with no physical data based on the query in your Declarative Pipelines

- TEMPORARY VIEW
- VIEW



Dataset Types Overview

TEMPORARY VIEW Overview

TEMPORARY VIEW

- Temporary views are only **persisted across the lifetime** of the pipeline and **private to the defining pipeline**
- They are **not registered** as an object to Unity Catalog
- Use the **CREATE TEMPORARY VIEW** statement
- TEMPORARY VIEWS are useful as **intermediate queries** that **are not exposed** to end users



Dataset Types Overview


VIEW Overview

VIEW

- Constructs a virtual table with no physical data based on the result-set of a SQL query in your pipelines
- They are **registered** as an object to Unity Catalog
- Use the **CREATE VIEW** statement


LIMITATIONS

- The pipeline must be a Unity Catalog pipeline
- Views cannot have streaming queries, or be used as a streaming source for a pipeline



Dataset Types Overview

Existing users of DLT will notice that we've evolved the names



The **semantics remain the same**, and **we'll support the old syntax** for compatibility.
Our goal is to simplify the syntax and match other systems



Dataset Types Overview

Pipeline dependencies are parsed automatically, code order is not required

Raw JSON → Bronze


```
CREATE OR REFRESH STREAMING TABLE  
1_bronze_db.orders_bronze AS  
SELECT  
*,  
current_timestamp() AS processing_time,  
_metadata.file_name AS source_file  
FROM STREAM read_files(  
"/Volumes/dbacademy/ops/labuser/orders",  
format => 'JSON');
```

Bronze → Silver

```
CREATE OR REFRESH STREAMING TABLE 2_silver_db.orders_silver  
AS  
SELECT  
order_id,  
timestamp(order_timestamp) AS order_timestamp,  
customer_id,  
notifications  
FROM STREAM 1_bronze_db.orders_bronze;
```

Silver → Materialized View

```
CREATE OR REFRESH MATERIALIZED VIEW 3_gold_db.gold_orders_by_date  
AS  
SELECT  
date(order_timestamp) AS order_date,  
count(*) AS total_daily_orders  
FROM 2_silver_db.orders_silver  
GROUP BY date(order_timestamp);
```





Lakeflow Declarative Pipeline Fundamentals

LECTURE

Simplified Pipeline Development



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Simplified Pipeline Development

Multi-file editor in Lakeflow Declarative Pipelines

What's New

1. Pipeline asset browser
2. Multi-file code editor with features for step-by-step development
3. Pipeline-specific toolbars
4. Interactive DAG
5. Data previews
6. Execution insights panels
 - Easier debugging
 - Faster validation with dry run

With Lakeflow Declarative Pipelines, you can write code using **.sql** or **.py** files (recommended), or use notebooks.

As of Data & AI Summit 2025

The screenshot shows the Lakeflow Declarative Pipelines interface. On the left, there's a sidebar with a tree view of pipeline assets (1). In the center, there's a code editor window (2) for a file named 'sample_trips_demo_pipeline.sql' containing SQL code for creating a materialized view. Above the code editor is a toolbar (3) with various icons. To the right of the code editor is a 'Pipeline graph' (4) showing two nodes connected by a flow. At the bottom, there's a 'Tables' tab (5) showing a list of tables with details like name, catalog, schema, type, and status. On the far right, there's a 'Pipeline settings' panel (6) with various configuration options like pipeline ID, type, name, and compute settings.





Lakeflow Declarative Pipeline Fundamentals

LECTURE

Common Pipeline Settings




Common Pipeline Settings

Overview


The screenshot shows the Databricks pipeline editor interface. On the left, there's a sidebar with a file tree and a central workspace for editing pipeline code. A red box highlights the gear icon in the sidebar, which is used to access pipeline settings. A large orange callout box contains the text: "You can edit Lakeflow Declarative Pipelines settings directly in the editor". An arrow points from this callout to the Pipeline settings sidebar on the right. The sidebar is also highlighted with a red border and contains the following information:

Pipeline settings	
Pipeline ID	ccb9eaa8-dfe2-4925-a36c-75e3c98b8a30
Pipeline type	ETL pipeline
Pipeline name	1 - Developing a Simple DLT Pipeline Project
Creator	peter.styliadis@databricks.com
Run as	peter.styliadis@databricks.com
Code assets	
Root folder	/Workspace/Users/p... DLT Pipeline Project
Source code	/Workspace/Users/p...s/orders_pipeline.sql
Default location for data assets	
Default catalog	dbacademy_peter_styliadis
Default schema	default
Compute	
Serverless	
Configuration	
source	/Volumes/dbacademy/ops/dbacademy_peter_styliadis
Tags	
Add tags	
Budget	

Common Pipeline Settings



Compute



Code Assets




Configuration
(parameters)



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Common Pipeline Settings



Compute



Code Assets



Configuration
(parameters)


SERVERLESS

- Databricks **recommends** developing new pipelines using serverless
- **Optimizes cost while maintaining** performance
- **Enables** you to focus on code rather than managing infrastructure
- Optional **Serverless Performance optimized** setting for time-sensitive workloads


- Serverless compute includes **Incremental refresh for materialized views**
- A cost-based optimizer to power fast and efficient transformations for materialized views on Serverless compute



Common Pipeline Settings



Compute



Code Assets



Configuration
(parameters)

SERVERLESS

- Databricks **recommends** developing new pipelines using serverless
- **Optimizes cost while maintaining** performance
- **Enables** you to focus on code rather than managing infrastructure
- Optional **Serverless Performance optimized** setting for time-sensitive workloads

CLASSIC (fixed size)

- Users need **permission** to create compute for Declarative Pipelines
- Workspace admins **can configure cluster policies** to provide users with access to compute resources for Declarative Pipelines

AUTOSCALING


- Enhanced autoscaling is **enabled by default** for all new pipelines.
- **Optimizes cluster utilization** by automatically allocating cluster resources based on workload volume



Common Pipeline Settings



Compute



Code Assets



Configuration
(parameters)

Pipeline settings

UI JSON X

Pipeline ID	5f9c6978-8d8b-4c21-b644-f861fb050656
Pipeline type	ETL pipeline
Pipeline name	New Pipeline 2025-04-10 10:31
Creator	peter.styliadis@databricks.com
Run as	peter.styliadis@databricks.com

Code assets

Root folder /Workspace/Users/peter/styliadis DLT Expectations

Source code /Workspace/Users/peter/orders_pipeline.sql

Configure paths

Default location for data assets

Default catalog dbacademy_peter_styliadis

Default schema default

Edit catalog and schema

← Pipeline source code

Pipeline root folder
Default location for source code. Automatically appended to Python sys.path to support imports of modules and packages in this folder from source code files.

/Workspace/Users/peter com/build-data-pipeline

Source code
SQL and Python files within specified folders or provided as individual file paths are processed when the pipeline runs.

/Workspace/Users/peter com/build-data-pipeline

+ Add path

The **Pipeline root folder** automatically includes all relevant files within the folder for the pipeline project (can be a Git folder)

The **source code** section references include subfolders (or files) for the pipeline (.py, .sql, notebooks)



Common Pipeline Settings



Compute



Code Assets



Configuration
(parameters)

A pipeline's configuration is **a map of key value pairs** that can be used to parameterize your code:

- Improve code readability/maintainability
- Reuse common parameters in multiple pipelines files

← Configuration

Key*	Value
source	/Volumes/dbacademy/ops/dbac

Add parameter

```
CREATE STREAMING TABLE bronze
SELECT *
FROM STREAM read_files(
    "${source}/orders",
    format => 'JSON');
```

In SQL reference the key's value using the \${key-name} syntax



Common Pipeline Settings

Common Settings

There are **various other settings** we will explore throughout the course, including:

- Pipeline settings
- Code assets
- Default location for data assets
- Compute
- Configuration
- Environment
- Tags
- Budget
- Advanced settings

The screenshot shows the 'Pipeline settings' page in the Databricks UI. The page is divided into several sections, each with a red border around its title. The sections include:

- Pipeline settings**: Shows Pipeline ID (ccb9eaa8-dfe2-4925-a36c-75e3c98b8a30), Pipeline type (ETL pipeline), Pipeline name (1 - Developing a Simple DLT Pipeline Project), Creator (peter.styliadis@databricks.com), and Run as (peter.styliadis@databricks.com).
- Code assets**: Shows Root folder (/Workspace/Users/p... DLT Pipeline Project) and Source code (/Workspace/Users/p...s/orders_pipeline.sql). There is also a 'Configure paths' button.
- Default location for data assets**: Shows Default catalog (dbacademy_peter_styliadis) and Default schema (default). There is a 'Edit catalog and schema' button.
- Compute**: Shows Serverless.
- Configuration**: Shows source (/Volumes/dbacademy/ops/dbacademy_peter_styliadis) and a 'Edit configuration' button.
- Tags**: Shows an 'Add tags' button.
- Budget**: Shows a blank section.





Lakeflow Declarative Pipeline Fundamentals

DEMONSTRATION

Developing a Simple Pipeline



Notebook: 2 - Developing a Simple Pipeline



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Lakeflow Declarative Pipeline Fundamentals

LECTURE

Ensure Data Quality with Expectations



Ensure Data Quality with Expectations

Creating Expectations

With Lakeflow Declarative Pipelines you can apply **data quality rules** to validate rows during ETL processes to gain insights into data integrity using the following

```
CONSTRAINT constraint-name EXPECT (column condition) [ON VIOLATION action]
```

Action	SQL Syntax Example	Result
WARN (default)	<code>CONSTRAINT valid_notification EXPECT (notifications IN ('Y', 'N'))</code>	<ul style="list-style-type: none">Invalid rows are still written to the targetLogs include counts of valid vs. invalid records and other metrics
DROP	<code>CONSTRAINT valid_date EXPECT (order_timestamp > "2021-01-01") ON VIOLATION DROP ROW</code>	<ul style="list-style-type: none">Invalid rows dropped from tableThe count is logged alongside other metrics
FAIL	<code>CONSTRAINT valid_id EXPECT (customer_id IS NOT NULL) ON VIOLATION FAIL UPDATE</code>	<ul style="list-style-type: none">Causes a failure of a single flow and does not cause other flows in your pipeline to fail.Manual intervention is required



Ensure Data Quality with Expectations

Adding Data Quality Expectations Example

```
CREATE OR REFRESH STREAMING TABLE 2_silver_db.orders_silver
(
    CONSTRAINT valid_notifications EXPECT (notifications IN ('Y', 'N')),
    CONSTRAINT valid_id EXPECT (customer_id IS NOT NULL) ON VIOLATION DROP ROW,
    CONSTRAINT valid_date EXPECT (order_timestamp > "2021-01-01") ON VIOLATION FAIL UPDATE
)
AS
SELECT
    Order_id,
    timestamp(order_timestamp) AS order_timestamp,
    Customer_id,
    notifications
FROM STREAM 1_bronze_db.orders_bronze;
```

WARN in metrics


DROP the row

FAIL THE PIPELINE



Ensure Data Quality with Expectations

Actions Overview



As of 2025Q2 materialized views that use **expectations** are **always fully refreshed**





Lakeflow Declarative Pipeline Fundamentals

DEMONSTRATION

Adding Data Quality Expectations



Notebook: 3 – Adding Data Quality Expectations



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Lakeflow Declarative Pipeline Fundamentals

LAB EXERCISE

Create a Pipeline



Notebook: 4 Lab - Create a Pipeline



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Building Lakeflow Declarative Pipelines



**Build Data Pipelines with Lakeflow Declarative
Pipelines**



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Section Learning Objectives

- Deploy a Lakeflow Declarative Pipeline in production by modifying configuration options like mode, schedule, email notifications and more.
- Analyze event logs and pipeline metrics to examine the entirety of a pipeline.
- Design and implement a Change Data Capture (CDC) Declarative Pipeline using APPLY CHANGES INTO to handle slowly changing dimensions (SCD).



Agenda

Section Overview – Building Pipelines

- **Streaming Joins Overview**
- **Deploying a Pipeline to Production**
- **Change Data Capture (CDC) Overview**
- **Change Data Capture with Apply CHANGE INTO**
- **Additional Features Overview**





Building Lakeflow Declarative Pipelines

LECTURE

Streaming Joins Overview



Streaming Joins Overview

Stream-Snapshot Join Overview

Streaming Table to Static Table




The goal is to **incrementally** join new data with a **static table**



Streaming Joins Overview

Stream-Snapshot Join Overview

Streaming Table to Static Table




The new data in the streaming table JOINS using all of the data in the lookup table



Streaming Joins Overview

Stream-Snapshot Join Overview

Streaming Table to Static Table



Only the new data in the streaming table JOINS using all of the data in the lookup table



Streaming Joins Overview

Joining Streaming Tables with a Materialized View

Streaming Table to Static Table

Streaming Table to Streaming Table with a Materialized View



The goal is to take **all rows from two streaming tables** and join them together each time the pipeline is run




Streaming Joins Overview

Joining Streaming Tables with a Materialized View

Streaming Table to Static Table

Streaming Table to Streaming Table with a Materialized View



The materialized view **efficiently computes all the data in the streaming tables and joins the rows.**




Streaming Joins Overview

Joining Streaming Tables with a Materialized View

Streaming Table to Static Table

Streaming Table to Streaming Table with a Materialized View



As new data is added to the streaming tables, the materialized view will again efficiently compute all the data in the streaming tables and joins the rows



Streaming Joins Overview

Stream-Stream Joins

Streaming Table to Static Table

Streaming Table to Streaming Table with a Materialized View

Outside the scope of the course

Incremental Stream to Stream Joins



The goal is to **incrementally** join new data from two tables, **past data is not used**

[Databricks Streaming and DLT](#) (Advanced)
Course





Building Lakeflow Declarative Pipelines


LECTURE

Deploying a Pipeline to Production



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Deploying a Pipeline to Production



- Enable pipeline **production** mode
- **Schedule** the pipeline
- Set up **email** notifications
- **Monitor** the pipeline with the **event log**



Deploying a Pipeline to Production

Common Production Tasks



Pipeline Modes



Scheduling



Email
Notifications



Monitor
pipelines



Deploying a Pipeline to Production

Common Production Tasks



Pipeline Modes



Scheduling



Email
Notifications



Monitor
pipelines

DEVELOPMENT Mode

- **Cluster Reuse** – Avoids overhead of restarts
- **Disables Pipeline Retries** – Immediate error detection and fixing
- **Optimized for Debugging** – Quick feedback loop for development

PRODUCTION Mode



Deploying a Pipeline to Production

Common Production Tasks



Pipeline Modes



Scheduling



Email
Notifications



Monitor
pipelines

DEVELOPMENT Mode

- **Cluster Reuse** – Avoids overhead of restarts
- **Disables Pipeline Retries** – Immediate error detection and fixing.
- **Optimized for Debugging** – Quick feedback loop for development.

PRODUCTION Mode

- **Cluster Restart** – Restarts for recoverable errors, including memory leaks and stale credentials.
- **Pipeline Retries** – Retries execution on specific errors (cluster startup failure).
- **Optimized for Reliability** – Ensures robustness and fault tolerance in production.



Deploying a Pipeline to Production

Common Production Tasks



Pipeline Modes



Scheduling



Email
Notifications



Monitor
pipelines

Triggered

- Can be triggered **manually** or run on a **schedule**
- **Refreshes** selected tables using data available at the start, then stops once the update is complete

Continuous

- **Continuously** processes new data to keep streaming tables and materialized views up to date in **near real time**
- Monitors dependencies and **updates only** when source data changes



Deploying a Pipeline to Production

Common Production Tasks



Pipeline Modes



Scheduling



Email
Notifications



Monitor
pipelines

You can configure **email notifications** when scheduling the pipeline

The screenshot shows a 'Schedule' dialog box. At the top, it says 'Every Day at 16 : 49'. Below that is a 'Timezone' dropdown set to '(UTC-04:00) Eastern Time (US and Canada)'. Under 'More options', there is a 'Notifications' section with a red border. It contains an input field with the email 'peter.fake@databricks.com' and three checked checkboxes: 'Start', 'Success', and 'Failure'. There is also a '+ Add' button and a trash icon. At the bottom right are 'Cancel' and 'Create' buttons.



Deploying a Pipeline to Production

Common Production Tasks



Pipeline Modes



Scheduling




Email
Notifications



Monitor
pipelines

- The **pipeline event log** captures all key information about the pipeline:
 - Audit logs
 - Data Quality Checks
 - Pipeline Progress
 - Data Lineage



Deploying a Pipeline to Production

Querying the Declarative Pipeline Event Log

- Publish the event log as a Delta table using the advanced settings
 - Specify the table location (catalog, schema)
 - Define the table name

Query the pipeline event log table

```
SELECT * FROM <catalog>.<schema>.<event_log_table_name>
```



Deploying a Pipeline to Production

Querying the Declarative Pipeline Event Log

- Publish the event log as a Delta table using the advanced settings
 - Specify the table location (catalog, schema)
 - Define the table name
- (DEFAULT) The Event Log is written as a **hidden Delta table**
 - Located in the pipeline **default catalog** and **schema**
 - View the [query the event log documentation](#) to access the hidden event log






Building Lakeflow Declarative Pipelines

DEMONSTRATION

Deploying Your Pipeline to Production



Notebook: 5 - Deploying Your Pipeline to Production



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Building Lakeflow Declarative Pipelines

LECTURE

Change Data Capture (CDC) Overview



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Change Data Capture (CDC) Overview

What is Change Data Capture?

- Change Data Capture (CDC) is a technique **used to track and capture changes in a data source** (such as a database, lakehouse or data warehouse).
- Then **applying those changes** into the Lakehouse or data warehouse

Source Data

INSERTS

UPDATES

DELETES

A Source table of changes,
currently this must be streaming
table

CDC is categorized into different types of **Slowly Changing Dimensions (SCD)**, which represent how **historical changes are handled**

Main types

SCD Type 1

SCD Type 2



Target Table

Table for changes to be applied



Change Data Capture (CDC) Overview

Slowly Changing Dimensions (SCD) Type 1 (1 of 2)

SCD Type 1 (Overwrite Target with Latest Values)

- When a record updates, the **previous record is simply overwritten** by its **key** with the new value.
- When a record is deleted by its **key**, the **record is removed**
- There is **no tracking of old keys (rows)**, only the **current data** is retained.

customers (*target table with all customers*)

CustomerID	Name	Address	ProcessDate
1	Peter	1 Blue Rd.	5/1/2025
2	Samarth	22 Front St	5/1/2025



*Update the **customers** table with the **new customer** information*


updates (*source table with updates*)

CustomerID	Name	Address	Operation	ProcessDate
1	Peter	123 Main St.	UPDATE	5/20/2025
2	null	null	DELETE	5/20/2025
3	Kostas	100 Athens Dr.	NEW	5/20/2025



Change Data Capture (CDC) Overview

Slowly Changing Dimensions (SCD) Type 1 (2 of 2)



Used SCD Type 1 when only the **current data matters** and historical values are irrelevant



Change Data Capture (CDC) Overview

Slowly Changing Dimensions (SCD) Type 2

SCD Type 2 (Historical Tracking/Versioning)

- When a record changes:
 - The **old record is preserved** with an additional column indicating its validity period (start date, end date, or a current flag) and a **new record is inserted** with the updated information.
 - When a **record is marked as deleted**, the record is kept and a column indicates the record is **inactive**.
- Used when **historical data is important**, and the system needs to track **how attributes change over time** (like tracking changes in a customer's address or status).

Example customers target table with SCD Type 2

CustomerID	Name	Address	ProcessDate	__START_AT	__END_AT	
1	Peter	123 Main St.	5/20/2025	5/20/2025	null	Null values indicate the current row
1	Peter	1 Blue Rd.	5/1/2025	5/1/2025	5/20/2025	Non null values indicate inactive (historic) rows
2	Samarth	22 Front St	5/1/2025	5/1/2025	5/20/2025	
3	Kostas	100 Athens Dr.	5/20/2025	5/20/2025	null	



Change Data Capture (CDC) Overview

Using `APPLY CHANGES INTO` to Perform SCD in Declarative Pipelines


customers (target table with all customers)

CustomerID	Name	Address	ProcessDate
1	Peter	1 Blue Rd.	5/1/2025
2	Samarth	22 Front St	5/1/2025




updates (source table with updates)

CustomerID	Name	Address	Operation	ProcessDate
1	Peter	123 Main St.	UPDATE	5/20/2025
2	Samarth	22 Front St	DELETE	5/20/2025
3	Kostas	100 Athens Dr.	NEW	5/20/2025



Change Data Capture (CDC) Overview





Building Lakeflow Declarative Pipelines

DEMONSTRATION

Change Data Capture with APPLY CHANGES INTO



Notebook: 6 - Change Data Capture with APPLY CHANGES INTO



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Building Lakeflow Declarative Pipelines

LAB EXERCISE

BONUS - APPLY CHANGES INTO with SCD Type 1

If you are interested in performing CDC with Declarative Pipelines we have provided an extra lab you can complete after class.



Notebook: 7 BONUS Lab - APPLY CHANGES INTO with SCD Type 1



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).



Building Lakeflow Declarative Pipelines

LECTURE

Additional Features Overview



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).

Additional Features Overview

What's Next: Features Outside This Course



FLows


Add **flows** for the **same target**, creating a union of the two data sources (or tables)



SINKS

Sink API **enables writing to destinations** outside of the pipeline:

- Kafka
- Delta
- Event hubs



Full DELTA Support

Improved Delta Support for MVs/STs

- Liquid clustering
- Row-level security and column masking
- Change data feed (CDF) from STs



Additional Features Overview

What's Next: Features Outside This Course



Full UC Support

- Publish to multiple catalogs & schemas
- Read **STs** and **MVs** in [Dedicated Access Mode](#)
- [Create a Unity Catalog pipeline by cloning a Hive metastore pipeline](#)



Better Performance

Improved performance throughout Lakeflow Declarative Pipelines!

- [Serverless](#)
- [Incremental refresh for materialized views](#)
- [Photon](#)



Databricks Asset Bundles (DABs)

[DABs](#) enable you to **programmatically validate**, **deploy**, and **run Databricks resources such as pipelines for CI/CD production workloads**

- [Automated Deployment with Databricks Asset Bundles](#) course



Summary and Next Steps



Course Learning Objective Recap

- Understand the core concepts and components of Lakeflow Declarative Pipelines, including the function and differences between streaming tables, materialized views, and temporary views.
- Identify and configure Lakeflow pipeline settings, such as compute, data assets, trigger modes, and advanced options.
- Develop a functional Lakeflow Declarative Pipeline using the new pipeline editor and SQL-based syntax.
- Incorporate data quality expectations into a pipeline to validate and enforce data integrity.
- Analyze event logs and pipeline metrics to understand the full execution and lifecycle of a Lakeflow Declarative Pipeline.
- Design and implement a Change Data Capture (CDC) to a pipeline using APPLY CHANGES INTO to handle slowly changing dimensions (SCD).



Next Steps

Additional resources for continuing the learning journey.

Data Engineering with Databricks

- Continue your learning through [self-paced](#) or [instructor-led](#) offerings
- The courses offer hands-on instruction in:
 - Databricks Data Science & Engineering Workspace
 - Databricks SQL
 - Declarative Pipelines
 - Databricks Repos
 - Databricks Task Orchestration
 - Unity Catalog

Data Engineer Associate Certification

- Validate your data and AI skills on Databricks by earning a Databricks credential
- [Exam information](#) and [exam guide](#)
- The exam covers:
 - Data Intelligence Platform
 - ELT With Spark SQL and Python
 - Incremental Data Processing
 - Production Pipelines
 - Data Governance





databricks



© Databricks 2025. All rights reserved. Apache, Apache Spark, Spark, the Spark Logo, Apache Iceberg, Iceberg, and the Apache Iceberg logo are trademarks of the [Apache Software Foundation](#).